

VA Diabetes Overview

Problem:

Veterans who receive care in Veterans Health Administration experience and suffer from a higher burden of obesity and diabetes compared to non-Veteran populations. The human and economic costs are staggering as reflected in the recent 2020 Obesity VA/DoD Clinical Practice Guideline Management of Adult Overweight and Obesity and the newly revised 2023 VA/DoD Clinical Practice Guideline Management of Type 2 Diabetes Mellitus. Most patients with diabetes mellitus are asymptomatic, which leads to delayed and more complex treatment.

Research Question:

Will empowering the veteran with knowledge and simple assessment promote healthier lifestyle choices that address and mitigate risk factors associated with T2DM?

Research Demographic - Veterans with BMI ≥ 25

Outcomes we hope to achieve over a 6 - 12 month follow up:

1. Weight loss: Evaluate percent of patients that achieved 5%; Evaluate percent of patients that achieved 10%
2. HbA1c reduction: Evaluate percent of patients that achieved HbA1c less than 6.5%; Evaluate percent of patients that reduced HbA1c by more than 1%

Objective:

The prevalence of Type 2 Diabetes Mellitus (T2DM) among veterans is a multifaceted issue. Environmental, genetic and lifestyle interactions have not been fully explored as predeterminants of T2DM. Glycated Hemoglobin (HbA1c) is one of the primary tools to diagnose T2DM. However, often times this is only monitored effectively after the onset of the disease. This research will develop a classification model to proactively identify veterans at high risk for T2DM, providing risk scores that enable earlier interventions and potentially prevent or delay disease development.

Imports

```
In [ ]: # Package imports
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import re
```

```
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')

# Set display options to show all rows
#pd.set_option('display.max_rows', 20)
```

```
In [ ]: import os

# Retrieve the 'Projects' environment variable
base_dir = os.environ.get('Projects', '/Users/ben/Projects')

# Construct the path to the Excel file
file_path = os.path.join(base_dir, 'VA_diab_review', 'update.xlsx')
```

```
In [ ]: # Loading xlsx file into a dataframe
df = (
    pd.read_excel(file_path, sheet_name = 'Sheet1')
    # Dropping columns with 0 non-null count
    .dropna(axis=1, how='all')
    # Lowercasing and replacing spaces with underscores and remove leading s
    .rename(columns=lambda x: x.strip().lower().replace(' ', '_').replace('-',
))
```

```
In [ ]: df.info()
df.tail()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1770 entries, 0 to 1769
Data columns (total 28 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   subject                                   574 non-null    float64
1   date_of_visit                             1181 non-null   object
2   age                                       1225 non-null   float64
3   sex                                       549 non-null    object
4   wt                                        1175 non-null   float64
5   sbp                                       1161 non-null   float64
6   dbp                                       1160 non-null   object
7   bmi                                       1105 non-null   object
8   obese                                    572 non-null    object
9   pre_dm                                   553 non-null    object
10  dm_t2                                    555 non-null    object
11  a1c                                       832 non-null    object
12  trig                                       679 non-null    float64
13  hdl                                       679 non-null    object
14  total_chol                             679 non-null    float64
15  ldl                                       675 non-null    object
16  medications_(gluc,_bp,_lipid,_statin)  502 non-null    object
17  cgm                                       21 non-null     object
18  mh_dx                                    330 non-null    object
19  wo_mh_dx                                217 non-null    object
20  metf                                      2 non-null      object
21  sglt2i                                   2 non-null      object
22  dpp4                                      3 non-null      object
23  su                                        2 non-null      object
24  ccb                                       3 non-null      object
25  ace/arb                                  4 non-null      object
26  diur                                      3 non-null      object
27  statin                                    4 non-null      object
dtypes: float64(6), object(22)
memory usage: 387.3+ KB
```

Out[]:

	subject	date_of_visit	age	sex	wt	sbp	dbp	bmi	obese	pre_dm	...
1765	NaN	2024-05-30 00:00:00	NaN	NaN	196.0	125.0	81	28.2	NaN	NaN	...
1766	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1767	573.0	2024-05-31 00:00:00	NaN	NaN	190.0	135.0	79	27.3	n	n	...
1768	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...
1769	574.0	2024-06-13 00:00:00	NaN	NaN	243.0	116.0	67	33	y	n	...

5 rows x 28 columns



Cleaning & Transformation

Research Audit

1. drop any nulls subsetting for ['date_of_visit', 'wt', 'a1c']
 2. create visit category for initial and follow up (address research questions)
 3. non-numeric character check necessary for expected numeric dtypes that imported as object
 4. laboratory nulls will not be imputed / outlier check (missing bmi impute by creating ht)
 5. a1c standard for diagnoses; create function to label encode for binary, multi-class / pre_dm and dm_t2 can be dropped
 6. create blood_pressure feature as ordinal EDA / potential drop
 7. create body_type instead of obese from bmi for EDA / potential drop
 8. drop med columns due to NULLS / drop age column due to error / drop sex due to balance
- HbA1c Target Labels

Healthy: Below 5.7%
 Prediabetes: 5.7% to 6.4%
 Diabetes: 6.5% or higher

blood_pressure

Normal: SBP is less than 120 and DBP is less than 80.
 Elevated: SBP is between 120 and 129 and DBP is less than 80.
 Hypertension Stage 1: SBP is between 130 and 139 or DBP is between 80 and 89.
 Hypertension Stage 2: SBP is 140 or higher or DBP is 90 or higher.
 Hypertensive Crisis: SBP is above 180 or DBP is above 120.
 UNDETERMINED: Any blood pressure reading that does not fit the above categories

body_type

Underweight: BMI less than 18.5
 Normal weight: BMI 18.5 to 24.9
 Overweight: BMI 25 to 29.9
 Obesity: BMI 30 and over

Pre-process and Model

8. date_of_visit, subject can be dropped before modeling
9. split dataframe into train/test
10. standarize and select classification models for eval

```
In [ ]: df = df.dropna(subset=['date_of_visit', 'wt', 'a1c']) # Drop rows with any
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 819 entries, 0 to 1769
```

```
Data columns (total 28 columns):
```

#	Column	Non-Null Count	Dtype
0	subject	548 non-null	float64
1	date_of_visit	819 non-null	object
2	age	721 non-null	float64
3	sex	524 non-null	object
4	wt	819 non-null	float64
5	sbp	814 non-null	float64
6	dbp	814 non-null	object
7	bmi	797 non-null	object
8	obese	547 non-null	object
9	pre_dm	540 non-null	object
10	dm_t2	541 non-null	object
11	a1c	819 non-null	object
12	trig	620 non-null	float64
13	hdl	620 non-null	object
14	total_chol	620 non-null	float64
15	ldl	616 non-null	object
16	medications_(gluc,_bp,_lipid,_statin)	429 non-null	object
17	cgm	19 non-null	object
18	mh_dx	313 non-null	object
19	wo_mh_dx	204 non-null	object
20	metf	2 non-null	object
21	sglt2i	2 non-null	object
22	dpp4	3 non-null	object
23	su	2 non-null	object
24	ccb	3 non-null	object
25	ace/arb	4 non-null	object
26	diur	3 non-null	object
27	statin	4 non-null	object

```
dtypes: float64(6), object(22)
```

```
memory usage: 185.6+ KB
```

```
In [ ]: # Address data types of a1c, hdl, and ldl by removing non-numeric characters
columns_of_interest = ['bmi', 'a1c', 'hdl', 'ldl']
unique_values = {col: df[col].unique() for col in columns_of_interest}

unique_values
```

```

Out[ ]: {'bmi': array([39.07, nan, 34.5, 31.67, 29.9, 29.59, 30.02, 24.7, 32.5, 38.
7,
34.12, 39.3, 30.5, 28, 42.9, 29, 33.8, 31.1, 41, 37.5, 37.7, 29.7,
33.1, 31.4, 30.7, 29.6, 26.7, 26, 45.8, 40.6, 55.3, 31.5, 28.6, 27,
33.09, 33.5, 43.1, 45.6, 41.5, 38.4, 36.2, 35.2, 36.6, 44.6, 46,
31.7, 33.4, 32.8, 35, 35.6, 34.6, 37.6, 30, 25, 42.7, 34.8, 32.1,
34.1, 38, 33.6, 30.2, 36, 36.9, 35.1, 38.2, 28.7, 35.8, 27.4, 50.5,
36.4, 36.8, 33.7, 30.3, 29.5, 28.8, 32.6, 40, 34.2, 30.8, 35.5,
26.8, 34.3, 25.5, 24.3, 27.7, 38.5, 32.4, 34.9, 31.2, 24, 23.5,
34.7, 41.1, 39.1, 35.3, 36.5, 31, 28.3, 40.5, 41.8, 40.2, 39.9,
31.8, 43.6, 33.3, 36.1, 46.5, 40.1, 39.7, 50.8, 35.9, 34.4, 45.3,
42.8, 36.7, 37.2, 32.3, 31.3, 32.9, 29.2, 27.9, 25.2, 28.9, 27.5,
27.6, 30.9, 33, 30.1, 57.5, 25.6, 30.6, 29.1, 35.4, 42.5, 29.3,
33.9, 31.6, 39.6, 47.5, 41.2, 37, 33.2, 30.4, 41.7, 41.6, 42, 28.4,
26.4, 36.3, 40.8, 43, 40.7, 41.9, 32, 45.5, 46.4, 23, 21, 23.6,
29.8, 25.1, 26.6, 25.7, 23.8, 39.2, '37.6q', 31.9, 45.9, 37.4,
26.2, 54, 24.4, 39, 40.3, '37?', 26.9, 24.6, 42.3, 32.2, 38.1,
43.4, 23.4, 42.4, 39.8, 37.1, 28.5, 43.3, 27.8, 28.2, 25.9, 27.3,
44, 47, 37.3, 28.1, 29.4, 37.9, 35.7, 37.8, 26.3, 32.7, 34, 38.8,
3, 27.2, 41.3, 43.8, 24.9, 42.1, 38.9, 38.3, 43.5, 55.6, 22.6,
49.3, 27.1, 54.7, 54.9, 23.2, 41.4, 48.4, 45.7, 44.4, 25.8, 47.1,
40.4, 42.6, 52.2, 25.3, 24.2, 44.8, 57.2, 38.6, 33.16],
dtype=object),
'a1c': array([6.4, 6.3, 5.8, 6.7, 5.3, 5.6, 9, 6.5, 6.1, 5.7, 6.9, 5.5, 5,
'5.7(2021)', 6, 5.4, 7.2, 5.1, 8.5, 7, 6.8, 5.9, 5.2, 6.6, 10.4,
9.2, 4.7, 8.6, 7.7, 7.9, 8.2, 4.4, 9.4, 4.9, 7.3, 7.5, 4.8, 7.1,
8.8, 11, 6.2, 9.5, 10.6, 10.3, 4.6, 8.3, 9.1, 8.1, 7.8, 8,
'7.1 (dex)', 4.5, 7.6, 10.9, '5.6 (5.8)', 9.7, '5.6 (5.7)',
'5.4 (5.7)', 9.8, 8.4, 10.7, '5.5 (5.7)', '8.5 (CGM 7.2)',
'6.6 (CGM)', '5.5(5.9)', '5.3(5.7)', 0.52, '5.3 (5.7)', 10, 10.5,
'5.6(5.8)', 9.9, 13.5, 11.9, 4, 14.1, 11.4, 7.4], dtype=object),
'hdl': array([54, nan, 52, 49, 33, 23, 28, 46, 27, 40, 36, 31, 34, 20, 50,
53,
44, 56, 68, 58, 55, 45, 29, 51, 35, 38, 41, 42, 32, 47, 39, 37,
'yy', 26, 30, 48, 43, 25, 76, 59, 61, 60, 57, 22, 62, 77, 73, 78,
69, 66, 63, 64, 75, 91, 71, 74, 18, 275, 101, 87, 24, 428, 98],
dtype=object),
'ldl': array([22.1, nan, 74, 84.6, 120.3, 53, 25, 130.9, 19, 96, 69, 116,
95,
100, 97, 108, 104, 115, 119, 137, 35, 152, 121, 101, 66, 157, 94,
165, 153, 144, 80, 164, 125, 149, 128, 129, 'nc', 113, 93, 61, 114,
33, 102, 134, 198, 175, 136, 130, 64, 23, 73, 213, 158, 172, 44,
162, 150, 29, 21, 87, 120, 109, 78, 138, 91, 65, 161, 103, 43, 51,
111, 156, 63, 107, 126, 90, 83, 173, 187, 123, 147, 140, 146, 88,
191, 151, 79, 76, 39, 77, 176, 181, 177, 99, 59, 31, 106, 131, 145,
81, 84, 194, 82, 143, 141, 124, 122, 110, 159, 118, 174, 127, 193,
155, 47, 42, 56, 75, 142, 112, 89, 0, 163, 62, 224, 105, 206, 86,
98, 168, 67, 132, 68, 50, 41, 117, 170, 133, 184, 72, 28, 221, 154,
135, 71, 186, 265, 234, 3, 92, 55, 7.2, 54, 85, 180, 232, 139, 27,
160, 273, 60, 208, 183, 167, 202, 207, 182, 169, 148, 220, 'NC',
200, 233, 197, 203, 32, 179, 188, 8, 58, 199, 166], dtype=object)}}

```

```

In [ ]: # Define the function to clean and convert bmi, a1c, ldl, hdl dtypes
def clean_and_convert_columns(df, columns):
    for column in columns:
        # Handle text within parentheses and any textual values for hdl and

```


```
df[column] = df[column].apply(lambda x: np.nan if isinstance(x, str)
# Remove anything within parentheses and convert to float
df[column] = df[column].astype(str).str.replace(r"\(.*\)", "", regex
return df

# Apply the function to clean and convert the columns
columns_to_clean = ['bmi', 'a1c', 'hdl', 'ldl']
clean_and_convert_columns(df, columns_to_clean)
```

Out []:

	subject	date_of_visit	age	sex	wt	sbp	dbp	bmi	obese	pre_dm	
0	1.0	2023-03-27 00:00:00	124.0	Male	264.0	129.0	78	39.07	YES	NO	
1	NaN	2023-06-08 00:00:00	124.0	NaN	258.0	116.0	62	NaN	NaN	NaN	
4	NaN	2024-05-16 00:00:00	NaN	NaN	233.0	104.0	79	34.50	NaN	NaN	
6	2.0	2023-03-28 00:00:00	124.0	Male	233.0	130.0	89	31.67	YES	NO	
7	NaN	2023-06-27 00:00:00	124.0	NaN	210.0	139.0	83	NaN	NaN	NaN	
...	
1758	NaN	2024-05-23 00:00:00	NaN	NaN	231.0	142.0	77	29.00	NaN	NaN	
1760	570.0	2024-05-31 00:00:00	NaN	NaN	200.0	132.0	75	32.80	y	n	
1762	571.0	2024-05-30 00:00:00	NaN	NaN	267.0	129.0	97	34.30	y	n	
1767	573.0	2024-05-31 00:00:00	NaN	NaN	190.0	135.0	79	27.30	n	n	
1769	574.0	2024-06-13 00:00:00	NaN	NaN	243.0	116.0	67	33.00	y	n	

819 rows x 28 columns



```
In [ ]: # convert date to datetime and forwardfill subject
df['date_of_visit'] = pd.to_datetime(df['date_of_visit'], errors='coerce')
df['subject'] = df['subject'].ffill()

In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 819 entries, 0 to 1769
```

```
Data columns (total 28 columns):
```

#	Column	Non-Null Count	Dtype
0	subject	819 non-null	float64
1	date_of_visit	819 non-null	datetime64[ns]
2	age	721 non-null	float64
3	sex	524 non-null	object
4	wt	819 non-null	float64
5	sbp	814 non-null	float64
6	dbp	814 non-null	object
7	bmi	795 non-null	float64
8	obese	547 non-null	object
9	pre_dm	540 non-null	object
10	dm_t2	541 non-null	object
11	a1c	802 non-null	float64
12	trig	620 non-null	float64
13	hdl	619 non-null	float64
14	total_chol	620 non-null	float64
15	ldl	574 non-null	float64
16	medications_(gluc,_bp,_lipid,_statin)	429 non-null	object
17	cgm	19 non-null	object
18	mh_dx	313 non-null	object
19	wo_mh_dx	204 non-null	object
20	metf	2 non-null	object
21	sglt2i	2 non-null	object
22	dpp4	3 non-null	object
23	su	2 non-null	object
24	ccb	3 non-null	object
25	ace/arb	4 non-null	object
26	diur	3 non-null	object
27	statin	4 non-null	object

```
dtypes: datetime64[ns](1), float64(10), object(17)
```

```
memory usage: 185.6+ KB
```

```
In [ ]: # all visits table
```

```
df_full = df.copy()
```

```
In [ ]: # Create new column for visit type (initial or followup)
```

```
df_sorted = df.sort_values(by=['subject', 'date_of_visit'])
```

```
# Function to label visits as 'initial' or 'follow_up', and drop middle visits
```

```
def label_visits(group):
```

```
    if len(group) > 2: # If there are more than 2 visits, drop the middle visits
```

```
        return pd.concat([group.head(1).assign(visit='initial'), group.tail(1)])
```

```
    elif len(group) == 2: # If there are exactly 2 visits
```

```
        return pd.concat([group.head(1).assign(visit='initial'), group.tail(1)])
```

```
    else: # If there is only 1 visit
```

```
        return group.assign(visit='initial')
```

```
# Apply the function to each group and re-assign to df
```

```
df = df_sorted.groupby('subject', as_index=False).apply(label_visits).reset_index()
```



```
In [ ]: df.columns
```

```
Out[ ]: Index(['subject', 'date_of_visit', 'age', 'sex', 'wt', 'sbp', 'dbp', 'bmi',
              'obese', 'pre_dm', 'dm_t2', 'a1c', 'trig', 'hdl', 'total_chol', 'ldl',
              'medications_(gluc,_bp,_lipid,_statin)', 'cgm', 'mh_dx', 'wo_mh_dx',
              'metf', 'sglt2i', 'dpp4', 'su', 'ccb', 'ace/arb', 'diur', 'statin',
              'visit'],
              dtype='object')
```

```
In [ ]: df.drop(columns=['age', 'sex', 'obese', 'medications_(gluc,_bp,_lipid,_statin)'])
```

```
In [ ]: # Function to create diabetes status
def categorize_diabetes_and_clean(df, a1c_column):
    """
    Categorizes HbA1c levels, adds binary and multi-class labels to the DataF
    and drops 'db_stat' and 'status' columns if they exist. Categories are t

    Parameters:
    - df: DataFrame containing a1c data.
    - a1c: The column name in the DataFrame that contains a1c values.

    The function modifies the DataFrame in-place and adds two new categorica
    - db_bin: Binary categorization (0 for Healthy, 1 for Diabetes).
    - db_multi: Multi-class categorization (0 for Healthy, 1 for Pre-Diabet

    Drops 'db_stat' and 'status' columns if present.
    """
    # Binary categorization
    db_bin_cats = pd.Categorical(df[a1c_column].apply(lambda x: 0 if x < 6.5
    df['db_bin'] = db_bin_cats.rename_categories(['Non-Diabetes', 'Diabetes'])

    # Multi-class categorization
    db_multi_cats = pd.Categorical(df[a1c_column].apply(lambda x: 0 if x < 5
    df['db_multi'] = db_multi_cats.rename_categories(['Healthy', 'Pre-Diabet

    # Drop 'db_stat' and 'status' columns if they exist
    df.drop(columns=['pre_dm', 'dm_t2'], errors='ignore', inplace=True)

    # Apply the function to create a new column for binary and multi-class categ
    categorize_diabetes_and_clean(df, 'a1c')
```

```
In [ ]: # Function for blood pressure categorization
def categorize_blood_pressure_and_clean(df, sbp_column, dbp_column):
    """
    Categorizes blood pressure readings into multi-class labels in the DataF
    and removes any redundant or irrelevant columns. The categories are trea

    Parameters:
    - df: DataFrame containing blood pressure data.
    - sbp: The column name in the DataFrame that contains systolic blood pre
    - dbp: The column name in the DataFrame that contains diastolic blood pr

    The function modifies the DataFrame in-place and adds a new categorical
    - blood_pressure: Multi-class categorization (Normal, Elevated, Hyperter
```

```

    Optionally drops any columns named 'old_bp' and 'previous_bp' if present
    """
    # Define multi-class categorization logic
    def multi_class_bp(sbp, dbp):
        if sbp < 120 and dbp < 80:
            return 'Normal'
        elif 120 <= sbp <= 129 and dbp < 80:
            return 'Elevated'
        elif (130 <= sbp <= 139) or (80 <= dbp <= 89):
            return 'Hypertension Stage 1'
        elif sbp >= 140 or dbp >= 90:
            return 'Hypertension Stage 2'
        elif sbp > 180 or dbp > 120:
            return 'Hypertensive Crisis'
        else:
            return 'Undetermined'

    # Apply categorization
    df['blood_pressure'] = pd.Categorical(df.apply(lambda row: multi_class_bp(
                                                categories=['Normal', 'Elevated',
                                                ordered=True)

    # Drop columns if they exist
    df.drop(columns=['old_bp', 'previous_bp'], errors='ignore', inplace=True)

categorize_blood_pressure_and_clean(df, 'sbp', 'dbp')

```

```

In [ ]: # Function for BMI categorization
def categorize_bmi_and_clean(df, bmi_column):
    """
    Categorizes BMI values into multi-class labels in the DataFrame,
    and removes any redundant or irrelevant columns. The categories are treated as ordered.

    Parameters:
    - df: DataFrame containing BMI data.
    - bmi_column: The column name in the DataFrame that contains BMI values.

    The function modifies the DataFrame in-place and adds a new categorical column.
    - body_type: Multi-class categorization (Underweight, Normal weight, Overweight, Obesity)

    Optionally drops any columns named 'old_bmi' and 'previous_bmi' if present
    """
    # Define multi-class categorization logic
    def multi_class_bmi(bmi):
        if bmi < 18.5:
            return 'Underweight'
        elif 18.5 <= bmi <= 24.9:
            return 'Normal weight'
        elif 25 <= bmi <= 29.9:
            return 'Overweight'
        else: # BMI of 30 and over
            return 'Obesity'

    # Apply categorization
    df['body_type'] = pd.Categorical(df[bmi_column].apply(multi_class_bmi),

```

```

categories=['Underweight', 'Normal weight', 'Overweight', 'Obesity']
ordered=True)

# Drop columns if they exist
df.drop(columns=['old_bmi', 'previous_bmi'], errors='ignore', inplace=True)

categorize_bmi_and_clean(df, 'bmi')

```

```

In [ ]: # mapping categorical values as ordinal for interpretability

db_multi_mapping = {'Healthy': 0, 'Pre-Diabetes': 1, 'Diabetes': 2}

db_bin_mapping = {'Non-Diabetes': 0, 'Diabetes': 1}

blood_pressure_mapping = {
    'Normal': 0, 'Elevated': 1, 'Hypertension Stage 1': 2,
    'Hypertension Stage 2': 3, 'Hypertensive Crisis': 4, 'Undetermined': 5
}

body_type_mapping = {
    'Underweight': 0, 'Normal weight': 1, 'Overweight': 2, 'Obesity': 3
}

# Apply the mappings
df['blood_pressure_c'] = df['blood_pressure'].map(blood_pressure_mapping)
df['body_type_c'] = df['body_type'].map(body_type_mapping)
df['db_multi_c'] = df['db_multi'].map(db_multi_mapping)
df['db_bin_c'] = df['db_bin'].map(db_bin_mapping)

# Convert these columns to numeric to ensure they are treated as such
df['blood_pressure_c'] = pd.to_numeric(df['blood_pressure_c'], errors='coerce')
df['body_type_c'] = pd.to_numeric(df['body_type_c'], errors='coerce')
df['db_multi_c'] = pd.to_numeric(df['db_multi_c'], errors='coerce')
df['db_bin_c'] = pd.to_numeric(df['db_bin_c'], errors='coerce')

```

```

In [ ]: df_follow = df[df['visit'] == 'follow_up']
df_init = df[df['visit'] == 'initial']
df_init.subject.nunique(), df_follow.subject.nunique()

```

Out[]: (547, 212)

```

In [ ]: # check how many subjects will be removed due to bmi < 25
df_init.query('bmi < 25').subject.nunique()

```

Out[]: 7

```

In [ ]: # create function for research questions

def create_research_df(df):
    # Filter for initial visits with BMI >= 25
    initial_visits = df[(df['bmi'] >= 25) & (df['visit'] == 'initial')]

    # Get the subject IDs of those with initial BMI >= 25
    subjects_with_high_bmi = initial_visits['subject'].unique()

    # Filter for follow-up visits for those subjects

```

```

follow_up_visits = df[(df['visit'] == 'follow_up') & (df['subject'].isin(
    subjects_with_follow_up))]

# Ensure that each subject with an initial visit also has a follow-up visit
subjects_with_follow_up = follow_up_visits['subject'].unique()

# Filter the initial visits to include only those subjects who have follow-up visits
initial_visits_with_follow_up = initial_visits[initial_visits['subject'].isin(
    subjects_with_follow_up)]

# Combine the initial visits and the selected follow-up visits
research_df = pd.concat([initial_visits_with_follow_up, follow_up_visits])

return research_df

# Apply the function to the original data
df_research = create_research_df(df)

```

```
In [ ]: df_research.value_counts('visit')
```

```
Out[ ]: visit
follow_up    206
initial      206
Name: count, dtype: int64
```

Exploratory Data Analysis

1. df_init = all initial visits
2. df_research = research questions on pre-post
3. df = use for feature importance / cor matrix / preprocess for modeling

```
In [ ]: df_init.drop(['date_of_visit', 'db_multi_c', 'db_bin_c'], axis=1).describe()
```

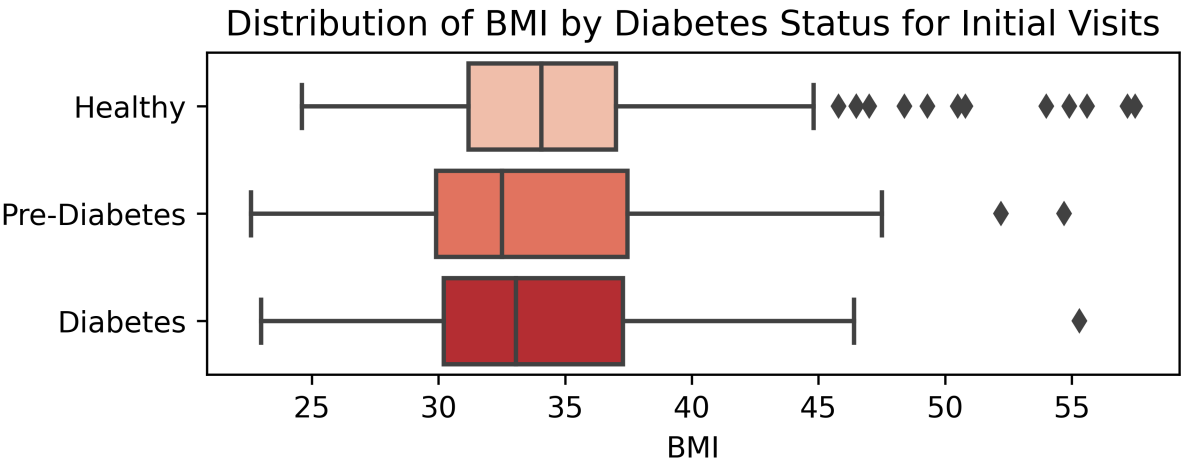
Out []:

	count	mean	std	min	25%	50%	75%	max
subject	547.0	282.334552	163.726419	1.00	141.5	282.0	421.5	574.0
wt	547.0	241.524680	40.885716	153.00	213.5	237.0	264.0	397.0
sbp	546.0	128.679487	13.184012	15.00	121.0	129.0	136.0	200.0
bmi	541.0	34.299815	5.565896	22.60	30.6	33.3	37.2	57.5
a1c	533.0	5.855197	1.194021	0.52	5.3	5.6	6.0	14.1
trig	531.0	199.990584	145.131411	25.00	107.0	159.0	243.5	1272.0
hdl	530.0	42.920755	19.847275	18.00	35.0	41.0	48.0	428.0
total_chol	531.0	197.531073	46.991734	52.00	165.0	197.0	224.0	419.0
ldl	488.0	116.963320	40.785945	0.00	91.0	116.5	143.0	273.0
blood_pressure_c	546.0	1.815018	0.863444	0.00	2.0	2.0	2.0	3.0
body_type_c	547.0	2.797075	0.433225	1.00	3.0	3.0	3.0	3.0

In []:

```
# Create the horizontal box plot using seaborn
plt.figure(figsize=(6, 2), dpi=600)
sns.boxplot(y="db_multi", x="bmi", data=df_init, orient="h", palette="Reds")

# Set the plot title and labels
plt.title("Distribution of BMI by Diabetes Status for Initial Visits")
plt.xlabel("BMI")
plt.ylabel("")
plt.show()
```




In []:

```
# Create the histogram with frequencies
plt.figure(figsize=(6, 4))
sns.histplot(data=df_init, x="db_multi", hue="body_type", multiple="fill", s

# Set the plot title and labels
plt.title("Distribution of Diabetes Categories by Body Type", fontsize=12)
plt.xlabel("")
```

```
plt.ylabel("Proportion")
plt.show()
```



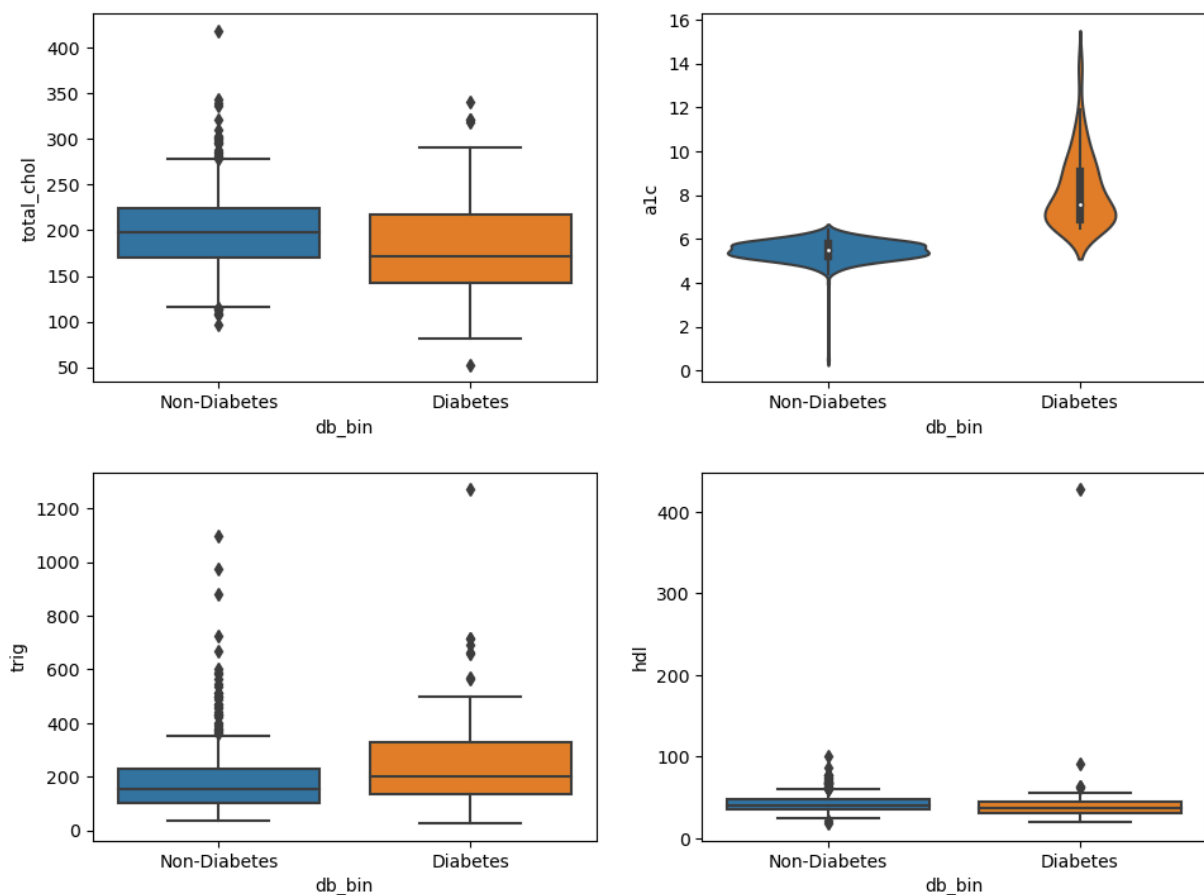
```
In [ ]: # Initial Visit Overview
```

```
fig, axes = plt.subplots(2, 2, figsize= (10, 8))

fig.suptitle('Labs by Patient Status (Visit = Initial)', fontsize=14)
sns.boxplot(ax=axes[0, 0], data=df_init, x='db_bin', y='total_chol', whis=1)
sns.violinplot(ax=axes[0, 1], data=df_init, x='db_bin', y='a1c', whis=1)
sns.boxplot(ax=axes[1, 0], data=df_init, x='db_bin', y='trig', whis=1)
sns.boxplot(ax=axes[1, 1], data=df_init, x='db_bin', y='hdl', whis=1)

plt.tight_layout(pad=2)
plt.show()
```


Labs by Patient Status (Visit = Initial)



```
In [ ]: # Create the count plot
sns.countplot(data=df_init, x='db_multi', hue='blood_pressure',
              palette='Reds', hue_order=['Normal', 'Elevated', 'Hypertension'])

# Set plot title and labels
plt.title('Clustered Barplot of Diabetes Status and Blood Pressure (visits = Initial)')
plt.legend(loc='upper right')
plt.xlabel('Diabetes Class')
plt.ylabel('Count')

# Show the plot
plt.show()
```



```
In [ ]: # Define numeric columns and subplot grid dimensions
numeric_cols = ['wt', 'sbp', 'dbp', 'bmi', 'a1c', 'trig', 'hdl', 'total_chol']
n_cols = 3
n_rows = (len(numeric_cols) + n_cols - 1) // n_cols

# Create figure and axes for subplots
fig, axes = plt.subplots(n_rows, n_cols, figsize=(5 * n_cols, 5 * n_rows))
fig.suptitle('Distribution of Numerics by Diabetes Status (Visit Type = Init)')

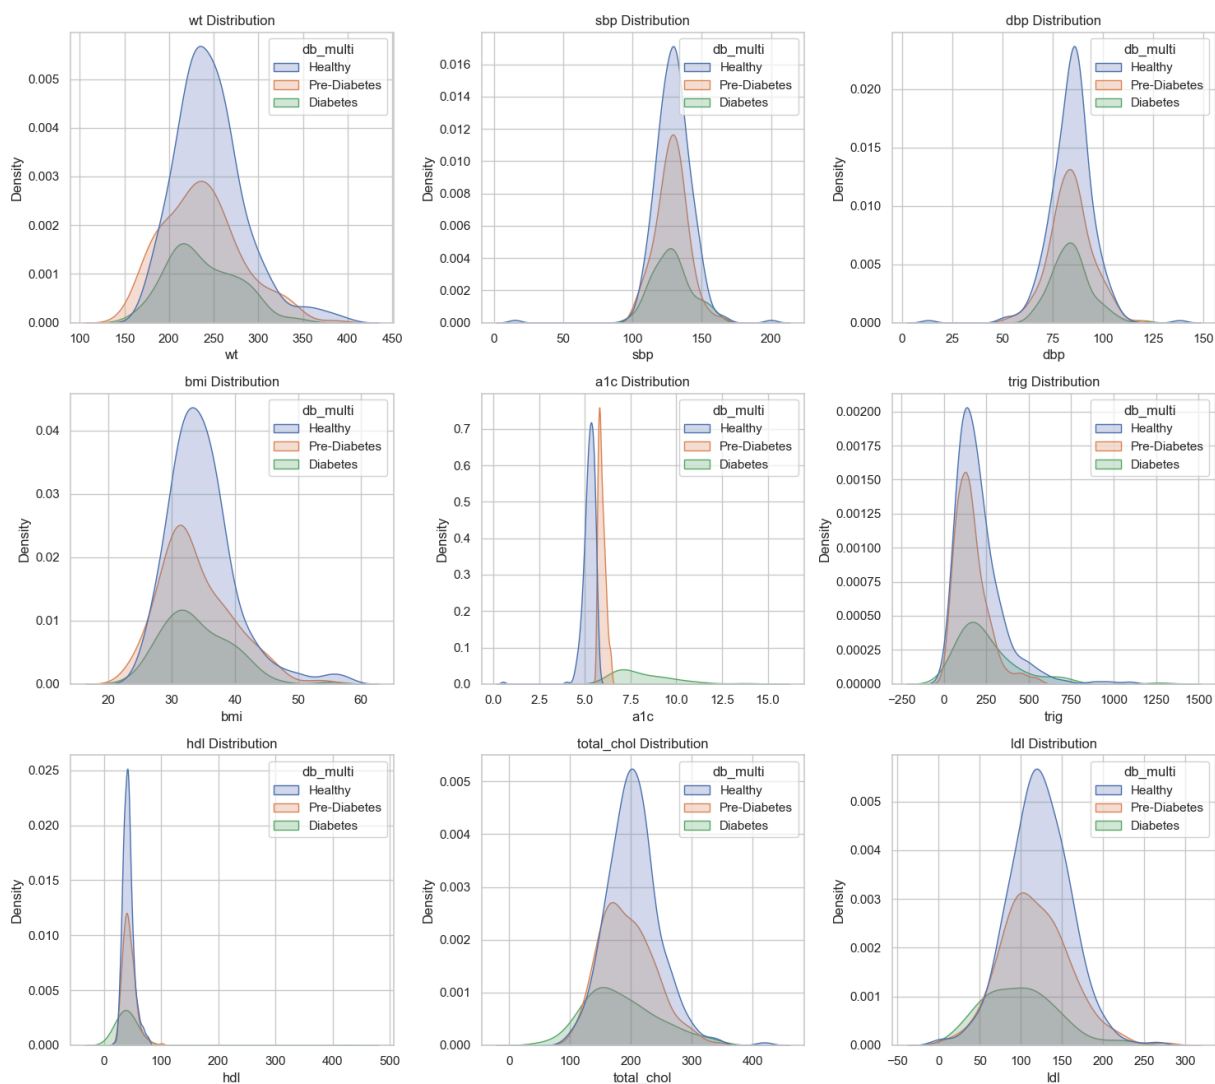
# Flatten axes array for easy iterating
axes = axes.flatten()

# Plot KDE for each numeric column with a hue
for i, col in enumerate(numeric_cols):
    sns.kdeplot(data=df_init, x=col, hue="db_multi", ax=axes[i], fill=True)
    axes[i].set_title(f'{col} Distribution')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('Density')

# Hide any unused axes if the number of numeric columns isn't a perfect mult
for ax in axes[len(numeric_cols):]:
    ax.set_visible(False)

plt.tight_layout(rect=[0, 0.03, 1, 0.95]) # Adjust layout to make room for
plt.show()
```



Distribution of Numerics by Diabetes Status (Visit Type = Initial)



```
In [ ]: # Select numeric columns for correlation matrix (including the newly coded c
numeric_cols = ['db_multi_c', 'body_type_c', 'blood_pressure_c', 'sbp', 'dbp',

# Calculate the correlation matrix
corr = df_init[numeric_cols].corr()

# Plot the heatmap of the correlation matrix using seaborn
plt.figure(figsize=(10, 8))
sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f", linewidths=.5)
plt.title('Correlation Matrix with Heatmap')
plt.show()
```



```
In [ ]: # Select only numeric columns from df_init
df_numeric = df_init.select_dtypes(include=[np.number])

# Remove 'db_multi_c' column if it exists in df_numeric to ensure it's not in
if 'db_multi_c' in df_numeric.columns:
    df_numeric = df_numeric.drop(columns=['db_multi_c', 'subject'])

# Compute the correlation matrix
correlation_matrix = df_numeric.corr()

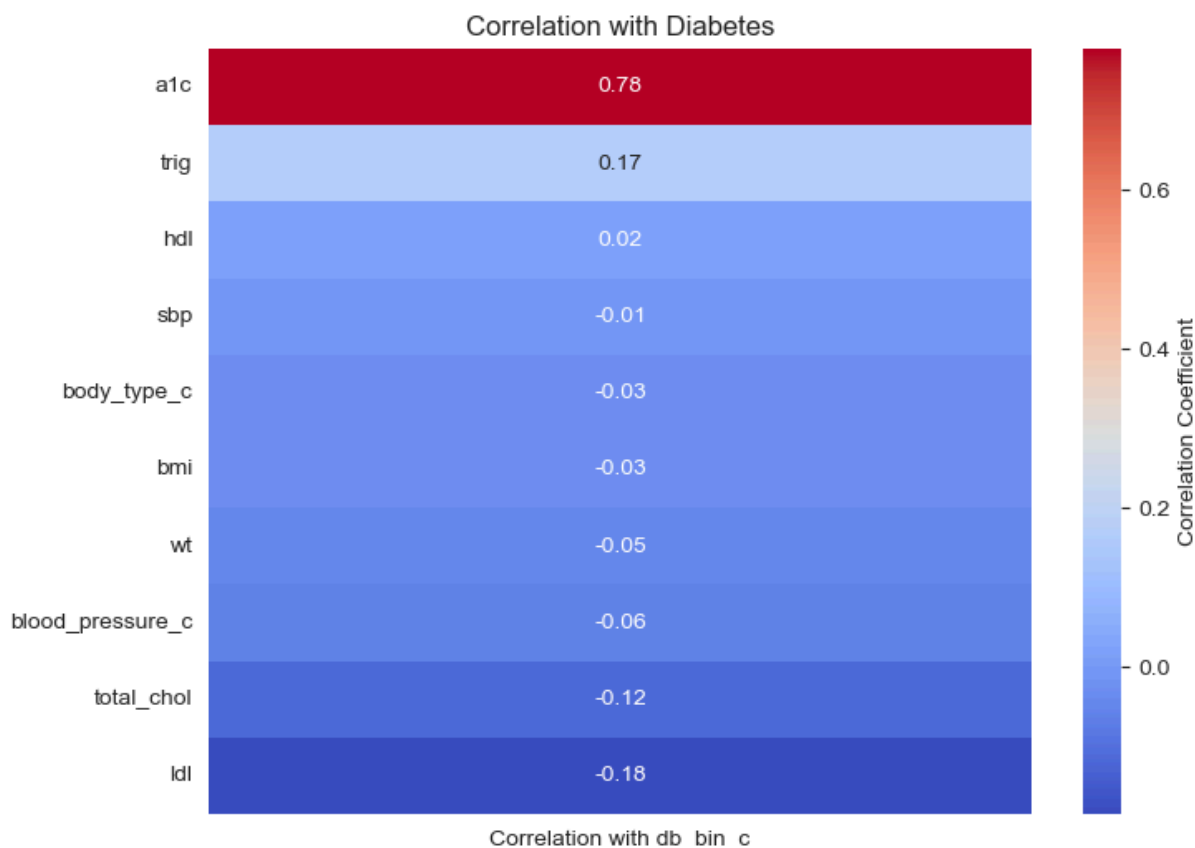
# Extract the correlation of 'db_bin_c' and drop its self-correlation
target_corr = correlation_matrix['db_bin_c'].drop('db_bin_c')

# Convert to DataFrame for heatmap compatibility and rename column for clarity
target_corr_df = target_corr.to_frame(name='Correlation with db_bin_c')

# Sort correlation values in descending order for better visualization
target_corr_sorted_df = target_corr_df.sort_values(by='Correlation with db_bin_c', ascending=False)

# Plotting setup
sns.set_style("white")
sns.set_palette("PuBuGn_d")
```

```
plt.figure(figsize=(8, 6)) # Adjust size to accommodate all variables if ne
sns.heatmap(target_corr_sorted_df, cmap="coolwarm", annot=True, fmt='.2f', c
plt.title('Correlation with Diabetes')
plt.show()
```



Research Outcomes

Outcomes we hope to achieve after intervention over a 6 - 12 month follow up

Weight loss: Evaluate percent of patients that achieved 5%; Evaluate percent of patients that achieved 10%

A1c reduction: Evaluate percent of patients that achieved HbA1c less than 6.5%;
Evaluate percent of patients that reduced HbA1c by more than 1%

```
In [ ]: # Define numeric columns and subplot grid dimensions
numeric_cols = ['sbp', 'dbp', 'bmi', 'a1c', 'trig', 'hdl', 'total_chol', 'ldl']
n_cols = 3
n_rows = (len(numeric_cols) + n_cols - 1) // n_cols

# Create figure and axes for subplots
fig, axes = plt.subplots(n_rows, n_cols, figsize=(5 * n_cols, 5 * n_rows))
fig.suptitle('Distribution of Labs by Between Visits', fontsize=20)

# Flatten axes array for easy iterating
axes = axes.flatten()

# Define a color palette
```

```

palette = "magma"

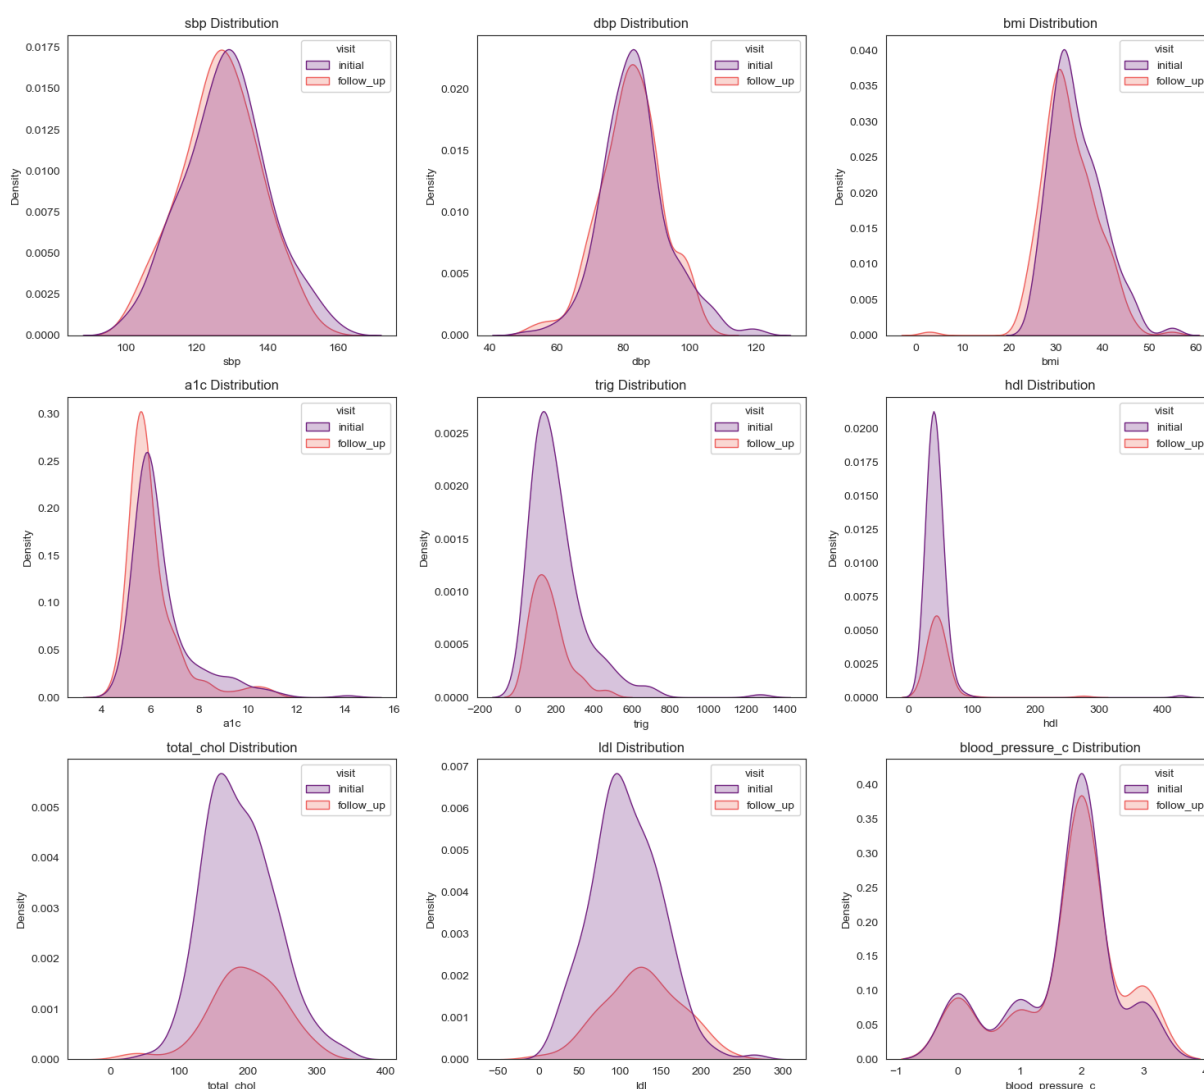
# Plot KDE for each numeric column with a hue
for i, col in enumerate(numeric_cols):
    sns.kdeplot(data=df_research, x=col, hue="visit", ax=axes[i], fill=True,
               axes[i].set_title(f'{col} Distribution')
               axes[i].set_xlabel(col)
               axes[i].set_ylabel('Density'))

# Hide any unused axes if the number of numeric columns isn't a perfect mult
for ax in axes[len(numeric_cols):]:
    ax.set_visible(False)

plt.tight_layout(rect=[0, 0.03, 1, 0.95]) # Adjust layout to make room for
plt.show()

```

Distribution of Labs by Between Visits



```

In [ ]: from scipy.stats import ttest_rel

# Perform the paired t-test between initial and follow-up bmi

# Filter DataFrame for initial and follow-up visits

```

```

df_initial = df_research[df_research['visit'] == 'initial'][['subject', 'bmi']]
df_follow_up = df_research[df_research['visit'] == 'follow_up'][['subject', 'bmi']]

# Merge the initial and follow-up DataFrames on subject
df_paired = pd.merge(df_initial, df_follow_up, on='subject', suffixes=('_ini', '_follow'))

# Drop any rows with NaN values in 'bmi'
df_paired.dropna(subset=['bmi_initial', 'bmi_follow_up'], inplace=True)

# Perform the paired t-test
t_stat, p_value = ttest_rel(df_paired['bmi_initial'], df_paired['bmi_follow_up'])

# Calculate Cohen's d
mean_diff = np.mean(df_paired['bmi_initial'] - df_paired['bmi_follow_up'])
std_diff = np.std(df_paired['bmi_initial'] - df_paired['bmi_follow_up'], ddof=1)
cohen_d = mean_diff / std_diff

t_stat, p_value, cohen_d

```

Out[]: (6.678290304440181, 2.4848629314932356e-10, 0.47824230201629464)

```

In [ ]: # Filter DataFrame for initial and follow-up visits for A1C similarly to BMI
df_initial_a1c = df_research[df_research['visit'] == 'initial'][['subject', 'a1c']]
df_follow_up_a1c = df_research[df_research['visit'] == 'follow_up'][['subject', 'a1c']]

# Merge the initial and follow-up A1C DataFrames on subject
df_paired_a1c = pd.merge(df_initial_a1c, df_follow_up_a1c, on='subject', suffixes=('_ini', '_follow'))

# Drop any rows with NaN values in 'a1c'
df_paired_a1c.dropna(subset=['a1c_initial', 'a1c_follow_up'], inplace=True)

# Merge the BMI and A1C change DataFrames to analyze together
df_changes = pd.merge(df_paired[['subject', 'bmi_initial', 'bmi_follow_up']], df_paired_a1c[['subject', 'a1c_initial', 'a1c_follow_up']], on='subject')

# Calculate changes in BMI and A1C
df_changes['bmi_change'] = df_changes['bmi_follow_up'] - df_changes['bmi_initial']
df_changes['a1c_change'] = df_changes['a1c_follow_up'] - df_changes['a1c_initial']

# Calculate the Pearson correlation between changes in BMI and changes in A1C
correlation_result = df_changes[['bmi_change', 'a1c_change']].corr().iloc[0, 1]
correlation_result

```

Out[]: 0.0909956793465215

```

In [ ]: # Perform the paired t-test for A1C between initial and follow-up visits
t_stat_a1c, p_value_a1c = ttest_rel(df_paired_a1c['a1c_initial'], df_paired_a1c['a1c_follow_up'])

# Calculate the mean difference and standard deviation for Cohen's d calculation
mean_diff_a1c = np.mean(df_paired_a1c['a1c_initial'] - df_paired_a1c['a1c_follow_up'])
std_diff_a1c = np.std(df_paired_a1c['a1c_initial'] - df_paired_a1c['a1c_follow_up'], ddof=1)
cohen_d_a1c = mean_diff_a1c / std_diff_a1c

t_stat_a1c, p_value_a1c, cohen_d_a1c

```

```
Out[ ]: (4.335726192019578, 2.2961707994308436e-05, 0.3050606303327793)
```

```
In [ ]: #calculating the percentage of patients achieving significant changes

# Filter and merge data for initial and follow-up visits directly
df_initial = df_research[df_research['visit'] == 'initial'][['subject', 'wt', 'a1c']]
df_follow_up = df_research[df_research['visit'] == 'follow_up'][['subject', 'wt', 'a1c']]
df_paired = pd.merge(df_initial, df_follow_up, on='subject')

# Calculate changes in weight and A1C
df_paired['wt_change'] = df_paired['wt_follow_up'] - df_paired['wt_initial']
df_paired['a1c_change'] = pd.to_numeric(df_paired['a1c_follow_up'], errors='coerce') - df_paired['a1c_initial']

# Determine significant weight loss and A1C reduction
df_paired['significant_wt_loss'] = (df_paired['wt_change'] / df_paired['wt_initial']) <= -0.05
df_paired['significant_10_wt_loss'] = (df_paired['wt_change'] / df_paired['wt_initial']) <= -0.1
df_paired['significant_a1c_reduction'] = df_paired['a1c_change'] <= -0.5
df_paired['significant_1_a1c_reduction'] = df_paired['a1c_change'] <= -1.0

# Calculate the percentage of patients achieving significant changes
percent_significant_wt_loss = 100 * df_paired['significant_wt_loss'].mean()
percent_significant_10_wt_loss = 100 * df_paired['significant_10_wt_loss'].mean()
percent_significant_a1c_reduction = 100 * df_paired['significant_a1c_reduction'].mean()
percent_significant_1_a1c_reduction = 100 * df_paired['significant_1_a1c_reduction'].mean()

(percent_significant_wt_loss, percent_significant_10_wt_loss, percent_significant_a1c_reduction, percent_significant_1_a1c_reduction)
```

```
Out[ ]: (41.262135922330096,
        13.106796116504855,
        32.038834951456316,
        13.592233009708737)
```

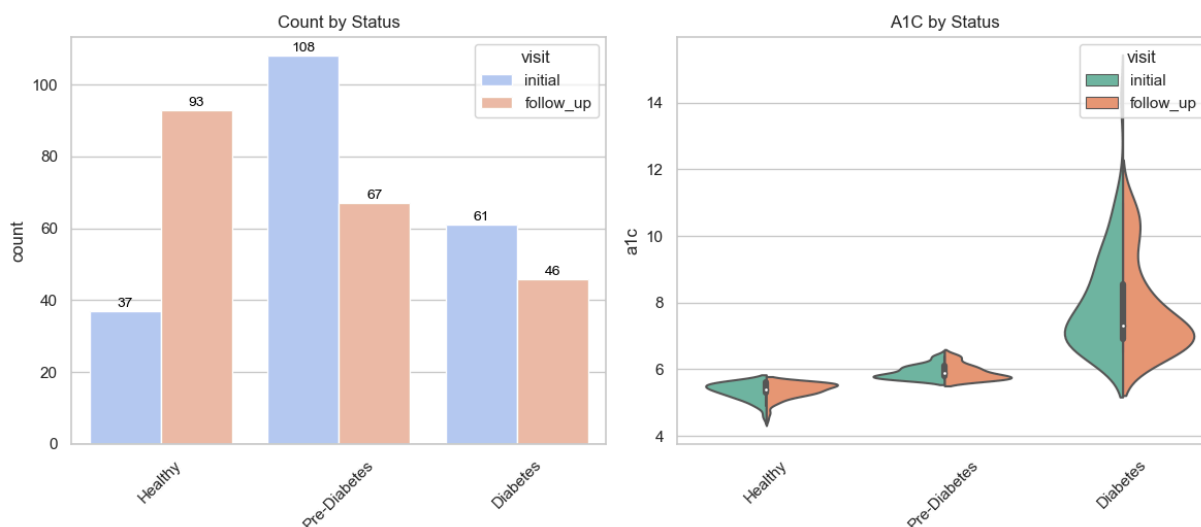
```
In [ ]: # Setup for the figure with two subplots
fig, axes = plt.subplots(1, 2, figsize=(12, 6))
fig.suptitle('Comprehensive Patient Status Analysis', fontsize=16)

# First subplot - Count plot for 'db_multi'
countplot1 = sns.countplot(ax=axes[0], data=df_research, x='db_multi', hue='visit')
axes[0].set_title('Count by Status')
axes[0].set_xlabel('db_multi')
axes[0].tick_params(axis='x', rotation=45)
for p in countplot1.patches:
    countplot1.annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2, p.get_height() + 5),
                        ha='center', va='bottom', fontsize=10, color='black',
                        textcoords='offset points')

# Second subplot - Violin plot for 'a1c' by 'db_multi'
sns.violinplot(ax=axes[1], data=df_research, x='db_multi', y='a1c', hue='visit')
axes[1].set_title('A1C by Status')
axes[1].set_xlabel('db_multi')
axes[1].tick_params(axis='x', rotation=45)

# Adjust layout
plt.tight_layout(rect=[0, 0, 1, 0.96]) # Adjust the layout to make room for the title
plt.show()
```

Comprehensive Patient Status Analysis

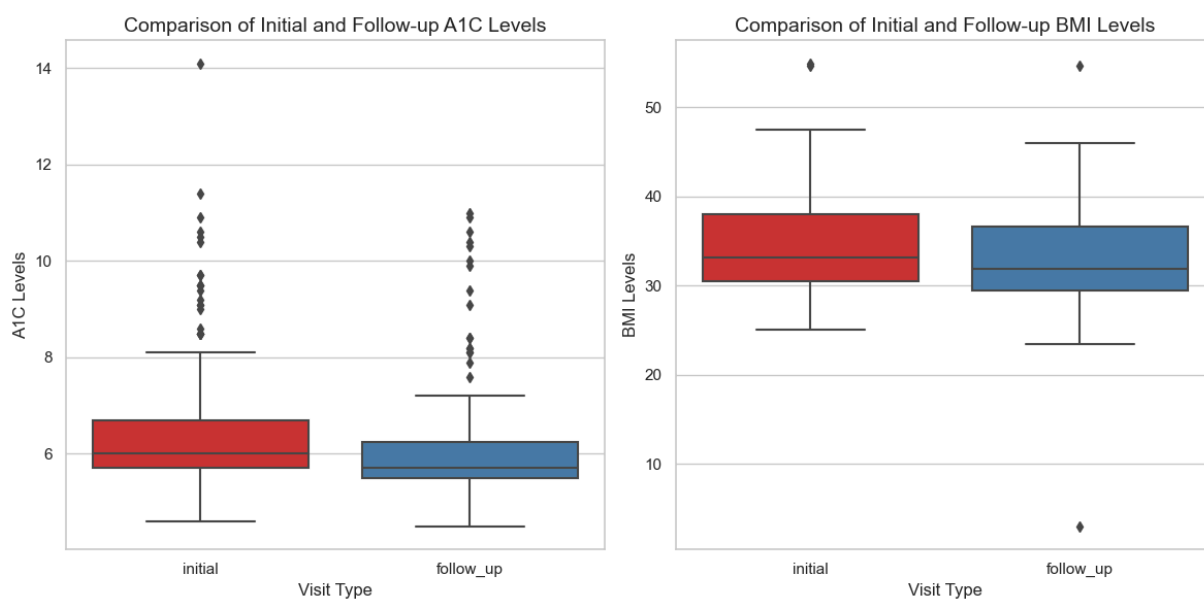


```
In [ ]: # Plotting the boxplot for A1C and BMI
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(12, 6)) # Use the same

# First subplot for A1C
sns.boxplot(x='visit', y='a1c', data=df_research, palette="Set1", ax=axes[0])
axes[0].set_title('Comparison of Initial and Follow-up A1C Levels', fontsize=12)
axes[0].set_xlabel('Visit Type')
axes[0].set_ylabel('A1C Levels')

# Second subplot for BMI
sns.boxplot(x='visit', y='bmi', data=df_research, palette="Set1", ax=axes[1])
axes[1].set_title('Comparison of Initial and Follow-up BMI Levels', fontsize=12)
axes[1].set_xlabel('Visit Type')
axes[1].set_ylabel('BMI Levels')

# Adjust layout for better spacing
plt.tight_layout()
plt.show()
```



```
In [ ]: import seaborn as sns
import matplotlib.pyplot as plt

# Set the aesthetic style of the plots
sns.set(style="whitegrid")

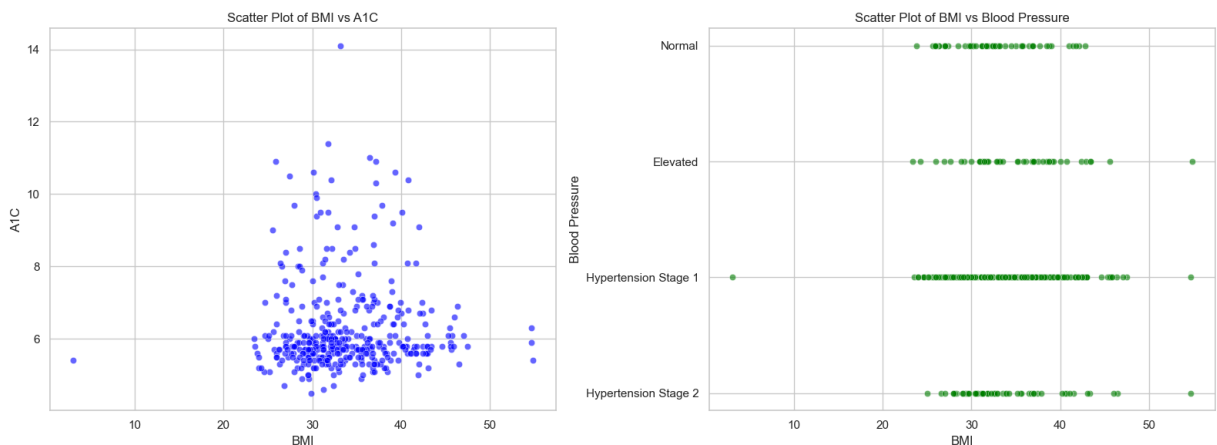
# Create a figure with subplots
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 6))

# Plot BMI vs A1C on the first subplot
sns.scatterplot(ax=axes[0], x='bmi', y='a1c', data=df_research, color='blue')
axes[0].set_title('Scatter Plot of BMI vs A1C')
axes[0].set_xlabel('BMI')
axes[0].set_ylabel('A1C')

# Plot BMI vs Blood Pressure on the second subplot
sns.scatterplot(ax=axes[1], x='bmi', y='blood_pressure', data=df_research, color='green')
axes[1].set_title('Scatter Plot of BMI vs Blood Pressure')
axes[1].set_xlabel('BMI')
axes[1].set_ylabel('Blood Pressure')

# Adjust layout
plt.tight_layout()

# Show the plots
plt.show()
```



Hypothesis

Empowering the veteran with knowledge and simple assessments allows clinicians to guide veterans in making scientifically based nutritional and behavioral changes, which will greatly improve their overall health and outcomes in the future.

Results BMI Reduction Analysis:

The paired t-test analysis comparing initial and follow-up BMI values of subjects revealed a significant difference, with a t-statistic of 6.678 and a p-value less than 0.05. This low p-value indicates that the observed difference is statistically significant, allowing us to reject the null hypothesis confidently. Additionally, Cohen's d, a measure

of effect size, was calculated to be 0.478, suggesting a medium effect size. This implies that the change in BMI over the observed period is not only statistically significant but also practically meaningful.

A1C Level Analysis:

The paired t-test analysis comparing initial and follow-up A1C values of subjects revealed a significant difference, with a t-statistic of 4.336 and a p-value less than 0.05. This low p-value indicates that the observed difference is statistically significant, allowing us to reject the null hypothesis confidently. Cohen's d, was calculated to be 0.305, suggesting a small to medium effect size.

Significant Weight Loss (at least 5%): Percentage of Patients: 41.26% This indicates that 41.26% of the patients achieved a weight loss of 5% or more between their initial and follow-up visits.

Significant Weight Loss (at least 10%): Percentage of Patients: 13.11% This indicates that 13.11% of the patients achieved a weight loss of 10% or more between their initial and follow-up visits.

Significant A1C Reduction (at least 0.5 units): Percentage of Patients: 32.04% This indicates that 32.04% of the patients achieved an A1C reduction of 0.5 units or more between their initial and follow-up visits.

Significant A1C Reduction (at least 1.0 unit): Percentage of Patients: 13.59% This indicates that 13.59% of the patients achieved an A1C reduction of 1.0 unit or more between their initial and follow-up visits.

These results highlight that a considerable portion of the patients experienced significant improvements in both weight and A1C levels over the observed period. Notably, over 40% of the patients achieved a meaningful weight loss of at least 5%, and approximately one-third saw a reduction of 0.5 units or more in their A1C levels. However, fewer patients achieved the more substantial targets of 10% weight loss and 1.0 unit A1C reduction.

Classification Model and Evaluation

```
In [ ]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV, cross_val_score, KFold
from sklearn.metrics import accuracy_score, classification_report, roc_curve

# List of columns to keep
```

```
columns = ['sbp', 'dbp', 'bmi', 'trig', 'hdl', 'total_chol', 'ldl', 'db_bin_c  
  
# Selecting the columns from the DataFrame  
df_mod = df[columns]
```

Regression Model and Evaluation